

■ AGILE INTEL

User Documentation

AI-Powered Risk Intelligence for Engineering Teams

Agile Intel synthesises data from Jira, Slack, and meeting transcripts into a single composite risk score and actionable coaching insights — so you see what is actually happening, not just what the board says.

1. What Does Agile Intel Do?

Most engineering teams live in Jira, but the real problems live in Slack threads, meeting notes, and tribal knowledge that never makes it onto the board. Agile Intel bridges that gap by continuously monitoring all three sources and surfacing the risk signals that matter.

- **Pulls data** from Jira, Slack, and Google Drive meeting transcripts on a schedule
- **Scores risk** across 7 independent signals — WIP aging, blockers, velocity, readiness, Slack signals, sentiment, and meeting alignment
- **Surfaces insights** — concrete blockers, decisions, risks, and scope changes extracted by AI from Slack messages and meeting transcripts
- **Answers questions** through a conversational AI coach that has live access to all of the above simultaneously

2. Key Concepts

Term	Meaning
Team	A Jira project mapped to one or more Slack channels. Each team has its own dashboard, risk score, and d
Risk Score	A composite 0–100 number representing overall sprint/delivery risk. Higher = more risk.
Signal	One of 7 risk dimensions (e.g. WIP Aging, Dependencies). Each is scored 0–100 independently.
Insight	A specific extracted piece of intelligence: a blocker, a decision, a risk, or a sentiment signal — with a sourc
Snapshot	A point-in-time capture of all 7 signal scores. Taken automatically on every sync cycle.
Sprint	Active, closed, and future sprints are all tracked. Velocity history is stored across all closed sprints.
Healthcheck	An on-demand deep analysis of your backlog quality across 11 criteria, scored issue-by-issue by AI.

3. Data Sources

Jira Cloud

Jira is the primary data source and is always required. Agile Intel reads all sprints (active, closed, future), issues, story points, statuses, priorities, assignees, issue links, and acceptance criteria. It tracks how long each issue has been in its current status for WIP aging calculations, and reads the full transition history of every in-progress issue.

Slack

Agile Intel reads messages from channels you configure (the Slack bot must be invited to each channel). Only the last 7 days of messages are analysed per sync. Claude AI processes each batch to extract risk signals: blockers, deployment failures, customer escalations, negative/positive sentiment, and unresolved threads.

Google Drive / Meeting Transcripts

Agile Intel syncs Google Docs that look like meeting transcripts by searching for files with names containing words like 'transcript', 'standup', 'retro', or 'planning'. Each document is exported as plain text and passed to Claude, which extracts blockers, decisions, action items, scope changes, and sentiment. Transcripts are

auto-matched to teams by detecting your Jira project key in the text. You can also upload transcripts manually by pasting raw text directly.

4. Pages & Features

4.1 Teams (Home)

The home screen shows all projects as cards. Each card displays: the project key and team name, the active sprint name, the current risk score with level label, a mini signal breakdown (WIP / DEPS / VEL / PBI / SLACK / SENT), the last sync timestamp, and quick-access buttons for Dashboard, Chat, and Timeline.

From the Teams page you can also:

- **SYNC ALL** — triggers a manual Jira sync for every team at once
- **+ NEW TEAM** — create a new team by entering a Jira project key and board ID

4.2 Dashboard

The main team view. Four panels are shown simultaneously:

Panel	What it shows
Risk Gauge	Circular dial with a needle showing composite risk score. Colour-coded LOW→MEDIUM→HIGH→CRITICAL
Signal Breakdown	Horizontal colour-coded bars for all 6 visible signals, each with its 0–100 score.
Risk Radar	Spider/radar chart plotting all 6 signals at once — useful for spotting the outlier dimension.
Composite Risk Trend	Line chart of the composite score over 14 days with reference lines at HIGH (50) and CRITICAL (75).

Header buttons:

- **HEALTHCHECK** — navigate to the Backlog Healthcheck report
- **TRANSCRIPTS** — navigate to the Transcripts page
- **ASK AI** — open the AI Chat coach
- **SYNC** — manually trigger a Jira sync + risk rescore for this team right now

4.3 AI Chat Coach

A multi-turn conversational interface powered by Claude. Every message you send is automatically enriched with live project context — you never need to copy-paste data into the chat. The AI knows everything the platform knows.

Context automatically injected per message:

Context Source	Detail
Risk snapshot	Current composite score + all 7 signals + velocity metrics
Active sprint	Sprint name, days remaining, goal, points done vs total
Full sprint board	Every issue grouped by status with story points, assignee, blocked flag, WIP aging flag
Sprint history	Last 4 closed sprints with committed/completed points and velocity
Recent insights	Up to 40 insights from the last 30 days, grouped by source
Meeting transcripts	Full text of the last 3 transcripts (up to 1,200 chars each)
Backlog healthcheck	Lead times, AI insights per group, top 5 worst items, recommended actions

Chat history is saved per session — the AI loads your last 10 turns as history so it remembers the conversation context.

4.4 Timeline

A chronological feed of all insights ever captured for a team. Useful for looking back at patterns, tracking when blockers first appeared, or reviewing what was discussed in a specific period. Insights can be filtered by source (Jira / Slack / Transcript), severity (low / medium / high / critical), and type (blocker / risk / decision / sentiment / dependency / scope change).

4.5 Backlog Healthcheck

An on-demand AI quality audit of your entire backlog. Unlike the sprint dashboard this analyses the last 6 months of issues — completed work, current sprint, and future backlog — across 11 quality criteria. Click **GENERATE REPORT** to run it. Analysis runs in the background (10–20 minutes) and the page polls automatically.

Report sections:

- **Lead Time Cards** — avg / story / bug / task lead times in days. Green <20d · amber 20–40d · red >40d
- **Backlog Depth Analysis Matrix** — 11 criteria × 3 groups (Completed Last Month / Planned This Month / Next Month Work List), showing % Good / Medium / Critical per cell as a stacked bar
- **AI Insights** — one 3-sentence paragraph per column identifying the key problems and priorities
- **Top 5 Items to Solve** — the five worst-scoring issues with AI-generated root cause explanations
- **Concrete Actions** — three prioritised team action items generated from the full analysis

The 11 quality criteria scored per issue:

Criterion	What it checks
Title Clarity	Specific, non-generic title that describes the actual work
Description Quality	Meaningful description with scope and context present
Acceptance Criteria	Clear, testable acceptance criteria are present
Estimation	Has story points or a size estimate assigned
Dependencies	External dependencies or blockers are identified
Customer Focus	User or customer value is clearly stated
Data-Driven	Data, metrics, or evidence backs the need
Vertical Slicing	Complete, end-to-end deliverable slice
Product Goal Alignment	Connected to a product goal or epic
OKR Alignment	Traceable to an OKR or strategic objective
Risk Profile	Risks, unknowns, or mitigations are mentioned

4.6 Transcripts

View and manage meeting transcripts. Each transcript card shows the title, date, participants, and the insights extracted from it.

Two ways to add transcripts:

- **Manual upload** — paste raw transcript text directly. Claude extracts insights immediately.
- **Google Drive sync** — connect your Google account once (OAuth). The system then automatically syncs relevant Docs from your Drive on the hourly schedule.

5. Risk Score Explained

The composite risk score (0–100) is a weighted average of 7 independent signals. Each signal is calculated from raw Jira, Slack, or transcript data.

Signal	Default Weight	What drives it up
WIP Aging	20%	Issues stuck In Progress beyond threshold (Stories 48h, Bugs 24h, Tasks 36h)
Dependencies	20%	Blocked issues — worse if the blocked issue is already In Progress
Velocity Trend	15%	Recent sprint velocity significantly below the prior average
PBI Readiness	15%	To-Do items missing story points, acceptance criteria, or assignee
Slack Signals	15%	Unresolved blockers, deployment failures, customer escalations from Slack
Team Sentiment	10%	Negative sentiment signals outnumbering positive ones across Slack + transcripts
Meeting Alignment	5%	Unresolved blockers, untracked action items, or scope changes from transcripts

Risk levels:

Score Range	Level	Meaning
0 – 24	LOW	Healthy. Monitor normally.
25 – 49	MEDIUM	Some signals worth watching.
50 – 74	HIGH	Active problems. Review blocked/aged items.
75 – 100	CRITICAL	Sprint delivery at serious risk. Immediate action needed.

Weights can be customised per team by an admin. For example, if Slack is not configured, increase the Jira-based weights to compensate.

6. How Automatic Syncs Work (Cron Jobs)

Agile Intel runs four background jobs on a continuous schedule to keep all data fresh. These jobs run silently — you never need to trigger them manually (though you can if you want fresher data immediately).

Job	Default Frequency	Configurable via	Requires
Jira Sync	Every 15 min	JIRA_SYNC_INTERVAL_MINUTES	Jira service account
Risk Scoring	Every 15 min	JIRA_SYNC_INTERVAL_MINUTES	Jira sync data
Slack Sync	Every 30 min	SLACK_SYNC_INTERVAL_MINUTES	Slack bot token
Transcript Sync	Every 60 min	TRANSCRIPT_SYNC_INTERVAL_MINUTES	Google OAuth connected

Job 1 — Jira Sync (every 15 minutes)

Loops through every active team and:

- Fetches all sprints (active, closed, future) for the team's Jira board
- Fetches all issues in the active sprint and all open issues in the project
- Updates sprint velocity, issue statuses, story points, WIP time, and transition history
- Calculates how long each issue has been in its current status
- Triggers Risk Scoring immediately after completing

Job 2 — Risk Scoring (every 15 minutes)

Runs immediately after each Jira sync completes. For each team:

- Reads the freshly synced Jira data from the database
- Computes all 7 signal scores
- Computes the weighted composite score
- Saves a new timestamped RiskSnapshot — this is what powers the trend chart
- Auto-generates insights for any high-scoring anomalies (e.g. an issue stuck for 72+ hours)

Each run produces a new snapshot, giving you a complete historical record of how your risk score has changed over time.

Job 3 — Slack Sync (every 30 minutes)

Only runs if a Slack bot token is configured. For each team with Slack channels set up:

- Fetches messages from the last 7 days from each configured channel
- Sends message batches to Claude, which extracts risk signals
- Saves any new insights to the database (deduplicates by content)
- Signal types extracted: blocker, deployment failure, customer escalation, help request, negative/positive sentiment, unresolved thread

Job 4 — Transcript Sync (every 60 minutes)

Only runs if at least one user has connected their Google account. For each user with a valid Google OAuth token:

- Searches Google Drive for recent documents with meeting-related names (contains 'transcript', 'standup', 'retro', 'planning', or 'Meet')
- Exports each document as plain text
- Passes text to Claude, which extracts blockers, decisions, action items, scope changes, and sentiment
- Auto-matches the transcript to the right team by looking for Jira project keys in the text
- Saves the transcript and its extracted insights to the database

Checking sync status

You can check which jobs are running and when they will run next via the API:

```
GET /api/v1/scheduler/status
```

7. Manual Sync & On-Demand Actions

You never have to wait for the scheduler — every sync can be triggered immediately from the UI:

Action	Where	What happens
■ SYNC button	Team Dashboard	Jira sync + risk rescore for that team right now
SYNC ALL button	Teams home page	Jira sync for all teams at once
■ GENERATE REPORT	Healthcheck page	Full backlog healthcheck analysis (runs in background, 10–20 min)
Manual transcript upload	Transcripts page	Processes transcript immediately; no wait time
Google Drive sync	Transcripts page	Triggers a one-off Drive scan for that team

8. User Roles

Role	Permissions
Admin	Full access: create teams, add/remove members, change risk weights, trigger any sync, access all teams
Coach	View all data, use AI chat, upload transcripts, trigger syncs — cannot create teams or manage users
Viewer	Read-only access to dashboards, timeline, and healthcheck reports

9. Integrations Setup

Jira Cloud (Required)

Jira is always required. You need a **service account** (a dedicated Jira user), its API token (generated in Jira → Account Settings → Security → API tokens), and your Jira Cloud base URL (e.g. <https://yourcompany.atlassian.net>). These are set as environment variables on the server. When creating a team you provide the **Jira Project Key** (e.g. ONB) and **Board ID** (visible in the Jira board URL).

Slack (Optional)

Create a Slack app with **channels:history** and **channels:read** scopes. Set the Bot Token (xoxb-...) as the `SLACK_BOT_TOKEN` environment variable. Then add channels to each team via the Teams page. The bot must be invited to each channel before it can read messages.

Google Drive / Transcripts (Optional)

Create a Google OAuth 2.0 client in Google Cloud Console with the Drive API enabled. Set `GOOGLE_CLIENT_ID` and `GOOGLE_CLIENT_SECRET` as environment variables. Each user then connects their own Google account via the Transcripts page. Agile Intel reads documents only — it never writes to your Drive.

10. Frequently Asked Questions

Q: How often does the risk score update?

A: Every 15 minutes automatically. You can force an immediate update using the ■ SYNC button on the Dashboard.

Q: Why is my Slack Signals score always 0?

A: Either Slack is not configured (no SLACK_BOT_TOKEN), no Slack channels have been added to your team, or no risk signals were found in the last sync window.

Q: How long does a Backlog Healthcheck take?

A: It depends on backlog size. Each batch of 20 issues takes about 15–20 seconds for Claude to score. A backlog of ~800 issues takes roughly 10–15 minutes. The page polls automatically and refreshes when complete.

Q: Does the AI chat see my latest Jira data?

A: Yes. Every message you send injects the latest data (from the last sync) as context, including the full sprint board, all insights from the last 30 days, recent transcripts, and the latest healthcheck report.

Q: Can I use Agile Intel without Slack?

A: Yes. The Slack Signals score will be 0 and those insights won't appear, but all Jira-based signals (WIP aging, dependencies, velocity, PBI readiness) work fully without Slack.

Q: What happens if Jira credentials expire?

A: The Jira sync job will fail. The last successful snapshot remains visible but won't update. The Dashboard shows the timestamp of the last sync so you can detect staleness.

Q: Can I adjust the risk score weights for my team?

A: Yes, as an Admin. Each team has configurable weights for all 7 signals (they must sum to 1.0). Use the Teams API (PATCH /api/v1/teams/{team_id}) or contact your admin.

Q: Are chat conversations stored?

A: Yes, per session. The AI loads your last 10 conversation turns as history so it remembers context within a session. History is stored in the database linked to your user ID.

Q: How is sentiment scored?

A: Each negative signal from Slack or transcripts adds points based on severity (critical=25, high=15, medium=8, low=4). Each positive signal subtracts 3 points. The score is capped at 100.